

Understanding the Transition from Memorization to Generalization in Diffusion Models for Cosmology

Jiaming Pan¹ Nicholas Kern¹ Dragan Huterer¹

¹Department of Physics, University of Michigan • jiamingp@umich.edu • 2026 AstroAI Workshop, CfA


M | **LSA** LEINWEBER INSTITUTE FOR THEORETICAL PHYSICS UNIVERSITY OF MICHIGAN



The question

Generative emulators promise **cheap surrogates** for expensive cosmological simulations — but an emulator that **memorizes** its training set can give a biased prior over fields. We train diffusion models on CAMELS HI maps and ask:

1. When does the model stop **copying** training fields and start **sampling** the underlying distribution?
2. Does conditional generation actually **respond to the requested cosmology** θ ?
3. Are generated fields **statistically faithful** (one-point PDF, power spectrum)?



If the learned prior collapses onto training islands, a downstream posterior can become biased even when the likelihood calculation is sharp.

Data & models


- **CAMELS** (IllustrisTNG) HI fields $\rightarrow 128 \times 128$ slices; 6 parameters $\Omega_m, \sigma_8, A_{SN1}, A_{AGN1}, A_{SN2}, A_{AGN2}$.
- **Diffusion**: UNet denoisers (widths 64 / 128 / 256); conditional variant injects θ at every level via cross-attention.
- **Sampling**: 50-step DPM-Solver, $8\times$ faster than the 500-step DDPM baseline at matched quality.
- **Training-set sweep**: $N = 2^6 \rightarrow 2^{15}$ fields, matched optimizer-update budget; runs on U-M Great Lakes HPC.

Detecting memorization

For each generated sample x_j , find its nearest training field in pixel, PCA, and SSCD embedding space:


$$s_j = \max_i \text{sim}(x_j, y_i), \quad \text{GL}(\tau) = 1 - \Pr[s_j > \tau]$$

with τ calibrated from leave-one-out train–train similarities. Independent models should reproduce the same distribution, not the same training images.



Calibration: does the generator obey θ ?


Ask the conditional model for **held-out samples**, then read the cosmology back from its generated maps with a **frozen encoder**. If the conditioning is used, changing θ_{ask} should move the recovered θ_{rec} rather than leaving it near the training-set average.



Encoder setup: VGG16 is frozen (ImageNet); only the regression head is trained on *real* slices from non-held-out simulations. Held-out test simulations are used only for encoder validation and generator calibration.

Memorization $\rightarrow$ generalization transition


Generated fields become less training-set-like with more data



- Score uses a **q95-calibrated threshold** ($\tau = 95$ th pct train–train NN similarity): higher means generated fields are less training-set-like in PCA space.
- **Small N_{2D}** : generated maps remain close to training fields.
- **Larger N_{2D}** : UNet-64 / 128 / 256 move toward novel samples; the wider UNet-256 transition is shifted to larger data.

Calibration results


Generated HI fields recover Ω_m



Main result: Ω_m recovery is much flatter under memorization (**slope 0.26**, $N_{2D} = 2^7$) than in the high-data regime (**slope 0.79**, $N_{2D} = 2^{14}$). Error bars are the 16–84 percentile spread over generated seeds at fixed θ_{ask} . We lead with Ω_m because the real-held-out VGG probe is strongest there ($R^2 = 0.91$; $\sigma_8 = 0.74$, feedback parameters weaker).

Physical fidelity

Generated power spectra approach the real reference during training



Across a single run, generated power spectra move toward the real reference from early to final checkpoints. Early checkpoints suppress power over many k bins; by the final checkpoint the generated $P(k)$ follows the real shape much more closely.

Takeaways

1. **Memorization is detectable.** Nearest-neighbor and reproducibility tests locate a data-size transition in CAMELS HI maps.
2. **Conditioning is recoverable.** High-data generated fields track θ_{ask} much better than memorization-regime fields.
3. **Fidelity requires generalization.** Useful emulators must be novel, statistically faithful, and calibrated before inference use.