Understanding the Transition from Memorization to Generalization in Diffusion Models for Cosmology

Jiaming Pan [1] Nicholas Kern [1] Dragan Huterer [1]

[1]Department of Physics, University of Michigan • jiamingp@umich.edu • 2026 AstroAI Workshop, CfA

LSA LEINWEBER INSTITUTE FOR THEORETICAL PHYSICS UNIVERSITY OF MICHIGAN

The question

Generative emulators promise **cheap surrogates** for expensive cosmological simulations — but an emulator that **memorizes** its training set can give a biased prior over fields. We train diffusion models on CAMELS HI maps and ask:

1. When does the model stop **copying** training fields and start **sampling** the underlying distribution?
2. Does conditional generation actually **respond to the requested cosmology** θ?
3. Are generated fields **statistically faithful** (one-point PDF, power spectrum)?

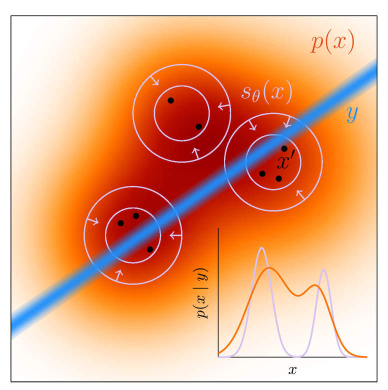

If the learned prior collapses onto training islands, a downstream posterior can become biased even when the likelihood calculation is sharp.

Data & models

- **CAMELS** (IllustrisTNG) HI fields $\to 128\times128$ slices; 6 parameters Ω_m, σ_8, A_{SN1}, A_{AGN1}, A_{SN2}, A_{AGN2}.
- **Diffusion**: UNet denoisers (widths 64 / 128 / 256); conditional variant injects θ at every level via cross-attention.
- **Sampling**: 50-step DPM-Solver, $8\times$ faster than the 500-step DDPM baseline at matched quality.
- **Training-set sweep**: $N = 2^6 \to 2^{15}$ fields, matched optimizer-update budget; runs on U-M Great Lakes HPC.

Detecting memorization

For each generated sample x_j, find its nearest training field in pixel, PCA, and SSCD embedding space:

$$s_j = \max_i \mathrm{sim}(x_j, y_i), \qquad \mathrm{GL}(\tau) = 1 - \Pr[s_j > \tau]$$

with τ calibrated from leave-one-out train–train similarities. Independent models should reproduce the same distribution, not the same training images.

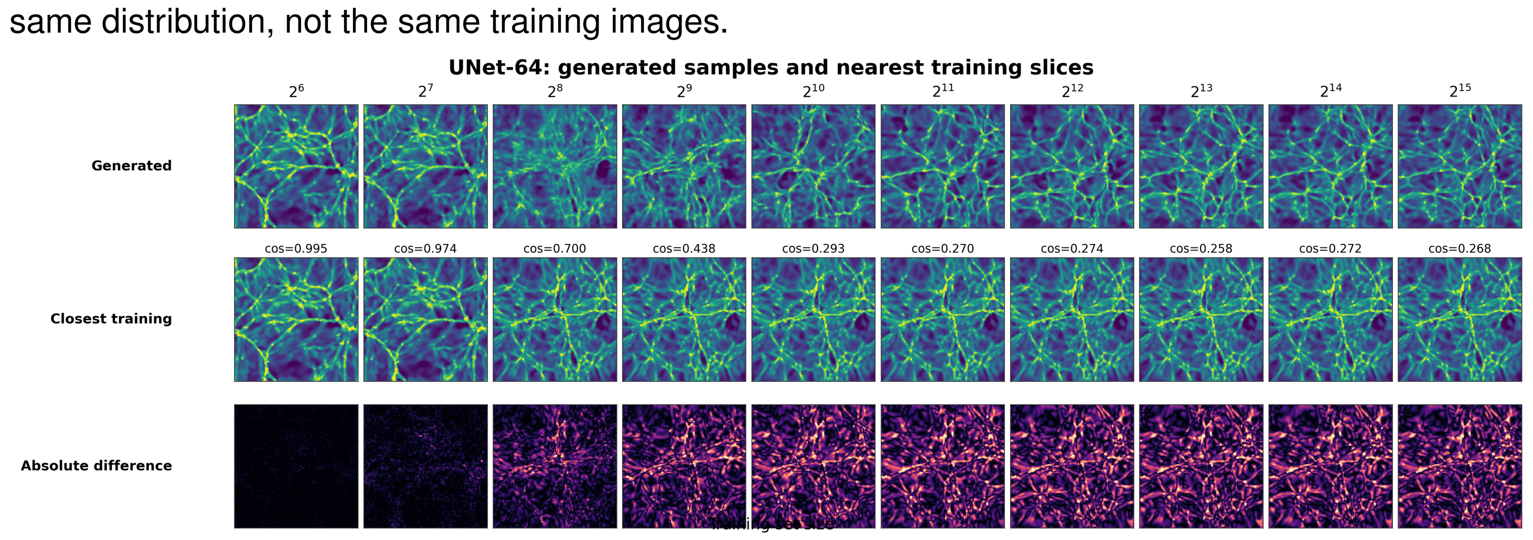

Calibration: does the generator obey θ?

Ask the conditional model for **held-out samples**, then read the cosmology back from its generated maps with a **frozen encoder**. If the conditioning is used, changing $\boldsymbol{\theta}_{\mathrm{ask}}$ should move the recovered $\boldsymbol{\theta}_{\mathrm{rec}}$ rather than leaving it near the training-set average.

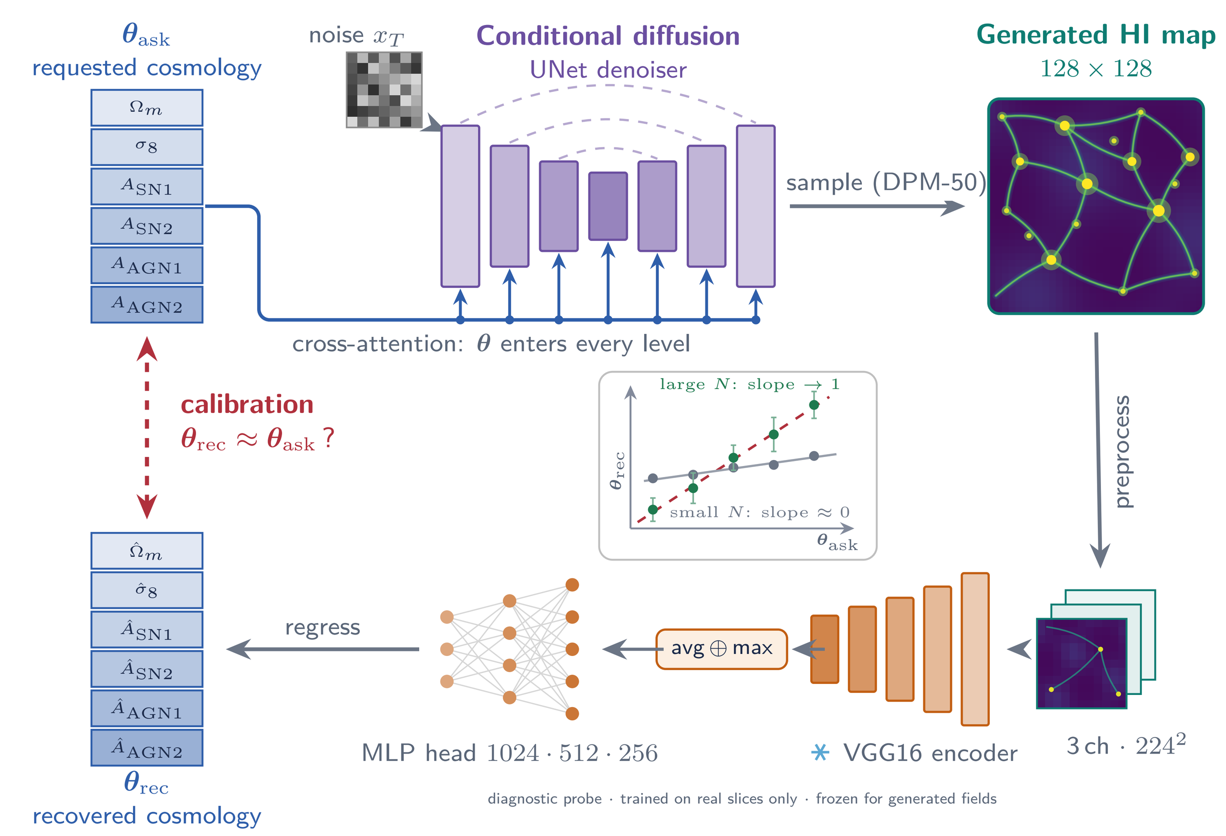

Encoder setup: VGG16 is frozen (ImageNet); only the regression head is trained on *real* slices from non-held-out simulations. Held-out test simulations are used only for encoder validation and generator calibration.

Memorization → generalization transition

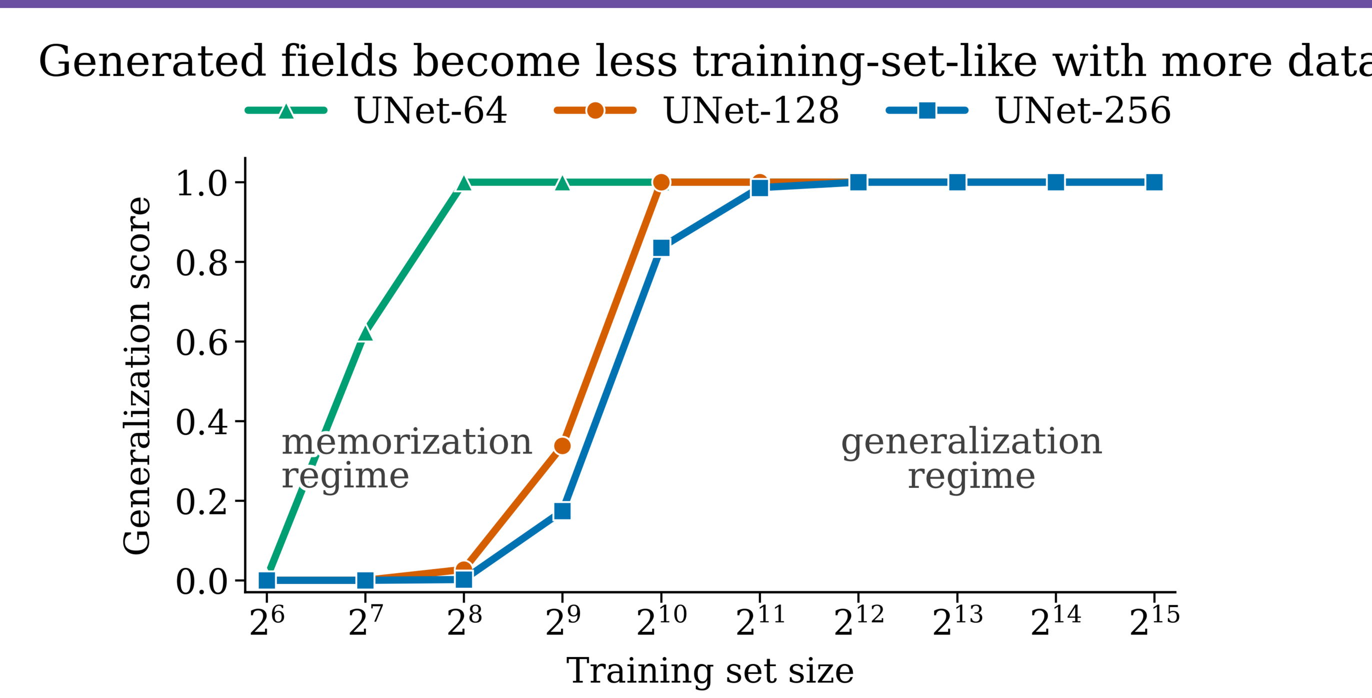

- Score uses a **q95-calibrated threshold** (τ = 95th pct train–train NN similarity): higher means generated fields are less training-set-like in PCA space.
- **Small** N_{2D}: generated maps remain close to training fields.
- **Larger** N_{2D}: UNet-64 / 128 / 256 move toward novel samples; the wider UNet-256 transition is shifted to larger data.

Calibration results

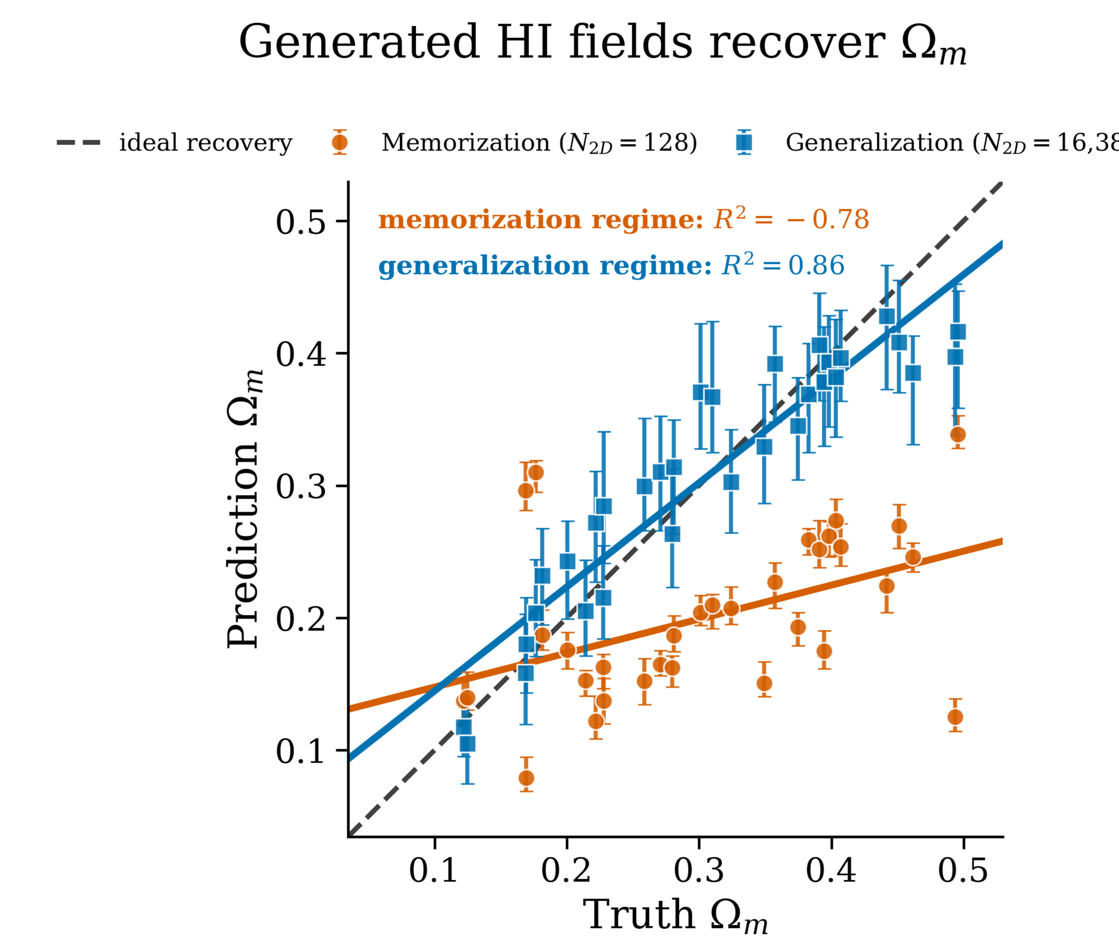

Main result: Ω_m recovery is much flatter under memorization (**slope 0.26**, $N_{2D} = 2^7$) than in the high-data regime (**slope 0.79**, $N_{2D} = 2^{14}$). Error bars are the 16–84 percentile spread over generated seeds at fixed θ_{ask}. We lead with Ω_m because the real-held-out VGG probe is strongest there ($R^2 = 0.91$; $\sigma_8 = 0.74$, feedback parameters weaker).

Physical fidelity

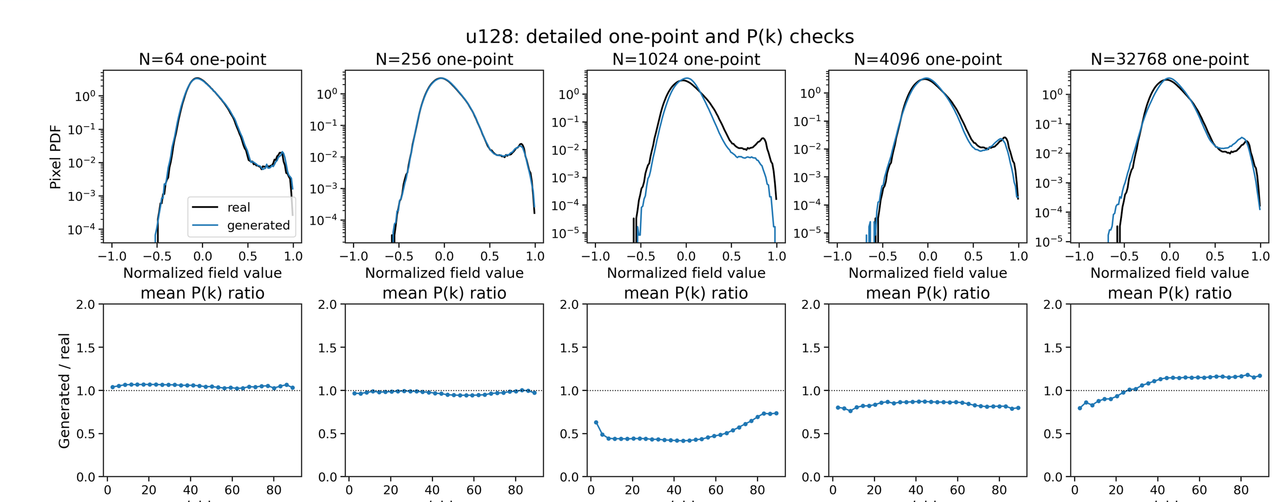

In high-data examples, the one-point PDFs overlap the real reference and the mean $P(k)$ ratio is mostly within $\sim 20\%$ over the resolved k range. Intermediate undertrained runs visibly fail this check, so novelty alone is not enough: generalization must also preserve field statistics.

Takeaways

1. **Memorization is detectable.** Nearest-neighbor and reproducibility tests locate a data-size transition in CAMELS HI maps.
2. **Conditioning is recoverable.** High-data generated fields track $\boldsymbol{\theta}_{\mathrm{ask}}$ much better than memorization-regime fields.
3. **Fidelity requires generalization.** Useful emulators must be novel, statistically faithful, and calibrated before inference use.