

Understanding the Transition from Memorization to Generalization in Diffusion Models for Cosmology

Jiaming Pan¹ Nicholas Kern¹ Dragan Huterer¹

¹Department of Physics, University of Michigan • jiamingp@umich.edu • 2026 AstroAI Workshop, CfA




The question

Generative emulators promise **cheap surrogates** for expensive cosmological simulations — but an emulator that **memorizes** its training set can give a biased prior over fields. We train diffusion models on CAMELS HI maps and ask:

1. When does the model stop **copying** training fields and start **sampling** the underlying distribution?
2. Does conditional generation actually **respond to the requested cosmology** θ ?
3. Are generated fields **statistically faithful** (one-point PDF, power spectrum)?

If the learned prior collapses onto training islands, a downstream posterior can become biased even when the likelihood calculation is sharp.



Data & models

- **CAMELS** (IllustrisTNG) HI fields $\rightarrow 128 \times 128$ slices; 6 parameters $\Omega_m, \sigma_8, A_{SN1}, A_{AGN1}, A_{SN2}, A_{AGN2}$.
- **Diffusion objective**: train a UNet denoiser with $x_t = \sqrt{\alpha_t}x_0 + \sqrt{1-\alpha_t}\epsilon$ and $v_t = \sqrt{\alpha_t}\epsilon - \sqrt{1-\alpha_t}x_0$; minimize $\mathcal{L}_{\text{diff}} = \mathbb{E} w(t) \|f_\theta(x_t, t, c) - v_t\|_2^2$; $c = \theta$ for conditional runs.
- **Sampling**: 50-step DPM-Solver, $8\times$ faster than the 500-step DDPM baseline at matched quality.
- **Training-set sweep**: $N = 2^6 \rightarrow 2^{15}$ fields, matched optimizer-update budget; runs on U-M Great Lakes HPC.


Detecting memorization

For each generated sample x_j , find its nearest training field in pixel, PCA, and SSCD embedding space:

$$s_j = \max_i \text{sim}(x_j, y_i), \quad \text{GL}(\tau) = 1 - \Pr[s_j > \tau]$$


with τ calibrated from leave-one-out train–train similarities. Independent models should reproduce the same distribution, not the same training images.

UNet-64: generated samples and nearest training slices



Calibration: does the generator obey θ ?


Ask the conditional model for **held-out samples**, then read the cosmology back from its generated maps with a **frozen encoder**. If the conditioning is used, changing θ_{ask} should move the recovered θ_{rec} rather than leaving it near the training-set average.



Encoder setup: VGG16 is frozen (ImageNet); only the regression head is trained on *real* slices from non-held-out simulations. Held-out test simulations are used only for encoder validation and generator calibration.

Memorization $\rightarrow$ generalization transition

Generated fields become less training-set-like with more data




- Score uses a **q95-calibrated threshold** ($\tau = 95\text{th}$ pct train–train NN similarity): higher means generated fields are less training-set-like in PCA space.
- **Small N_{2D}** : generated maps remain close to training fields.
- **Larger N_{2D}** : UNet-64/128/256 move toward novel samples; the wider UNet-256 transition is shifted to larger data.

Calibration results

Generated HI fields recover Ω_m


-- ideal recovery • mem. regime ($N_{2D} = 2^7$), slope=0.26 ■ gen. regime ($N_{2D} = 2^{14}$), slope=0.79



Main result: error bars show 16–84% seed scatter at fixed θ_{ask} . In the high-data regime they bracket the requested Ω_m ; under memorization they can be small but centered on the wrong value. Slopes (0.26 vs 0.79) are secondary summaries; Ω_m is shown because the held-out VGG encoder is strongest there ($R^2 = 0.91$).

Physical fidelity

Generated power spectra approach the real reference during training



Across a single run, generated power spectra move toward the real reference from early to final checkpoints. Early checkpoints suppress power over many k bins; by the final checkpoint the generated $P(k)$ follows the real shape much more closely.

Takeaways

1. **Memorization is detectable.** Nearest-neighbor and reproducibility tests locate a data-size transition in CAMELS HI maps.
2. **Conditioning is recoverable.** High-data generated fields track θ_{ask} much better than memorization-regime fields.
3. **Fidelity requires generalization.** Useful emulators must be novel, statistically faithful, and calibrated before inference use.